
Apache Spark

Introduction

BIG DATA

ANALYTICS

STORAGE

TECHNOLOGIES

SOFTWARE

USE

ONE

SINCE

REQUIRING

BUSINESS

MOVING

UBIQUITOUS

SAN

RADIO-FREQUENCY

SOLID

WIRELESS

TOLERABLE

SIZE

NEEDED

QUALITIES

PETABYTES

INTERNET

LOGS

PROCESSING

SHARED

DISK

TIME

RELATIONAL

SYSTEMS

NETWORKS

INFORMATION

RECORDS

COMPLEX

DATABASES

SETS

EXAMPLES

APPLIED

ELAPSED

CURRENT

THOUGHT

AMOUNT

SENSOR

ARCHIVES

ABILITY

TARGET

DIFFICULTY

INDEXING

CITATION

COST

CONTINUES

TERABYTES

MPP

RESEARCH

WOLFS

PRESENTATIONS

FORMS

CAPACITY

TENS

NEW

PRACTITIONERS

PC

CURRENTLY

DESKTOP

CREATED

WORKING

TYPES

USING

MASSIVELY

GARTNER

RELATED

BIOLOGICAL

HUNDREDS

ALSO

PERFORMANCE

ZETTABYTES

GENOMICS

SET

COMPUTING

TOOLS

WITHIN

PROCESS

DEFINITION

SEARCH

RECONSIDER

OPPORTUNITIES

CONNECTOMICS

USED

DISTRIBUTED

MAY

MANAGEMENT

CAPTURE

MANAGE

GROW

STORE

DEFINING

CASE

EVERY

LARGER

How to define Big Data?

- Volume
- Velocity
- Variety
- Variability
- Complexity



Overview

- developed by Mateu Zaharia as a part of PhD thesis at UC Berkley 2013
- written in java
- distributed querying and processing engine
- up to 100 times faster than Apache Hadoop when data is stored in memory and up to 10 times when accessing disk

The logo for Spark SQL, featuring the word "Spark" in a black, sans-serif font with an orange star above the "a", followed by "SQL" in a larger, bold, black, sans-serif font.The logo for Spark MLlib, featuring the word "Spark" in a black, sans-serif font with an orange star above the "a", followed by "MLlib" in a smaller, black, sans-serif font.The logo for Spark Streaming, featuring the word "Spark" in a black, sans-serif font with an orange star above the "a", followed by "Streaming" in a smaller, black, sans-serif font.The logo for GraphX, featuring a stylized graph icon with nodes and edges, followed by the word "GraphX" in a black, sans-serif font.The Apache Spark logo, featuring the word "APACHE" in a small, black, sans-serif font above the word "Spark" in a large, black, sans-serif font, with an orange star above the "a".The Hadoop HDFS logo, featuring a yellow elephant icon to the left of the word "hadoop" in a blue, sans-serif font, with "HDFS" in a yellow, sans-serif font below it.The Amazon S3 logo, featuring the Amazon logo (a stylized "A" made of orange and yellow cubes) to the left of the word "amazon" in a black, sans-serif font, with "web services" in a smaller, black, sans-serif font below it, and "S3" in a large, bold, black, sans-serif font to the right.



Spark ecosystem

- Spark SQL
- Spark MLlib
- Spark Streaming
- Graphx

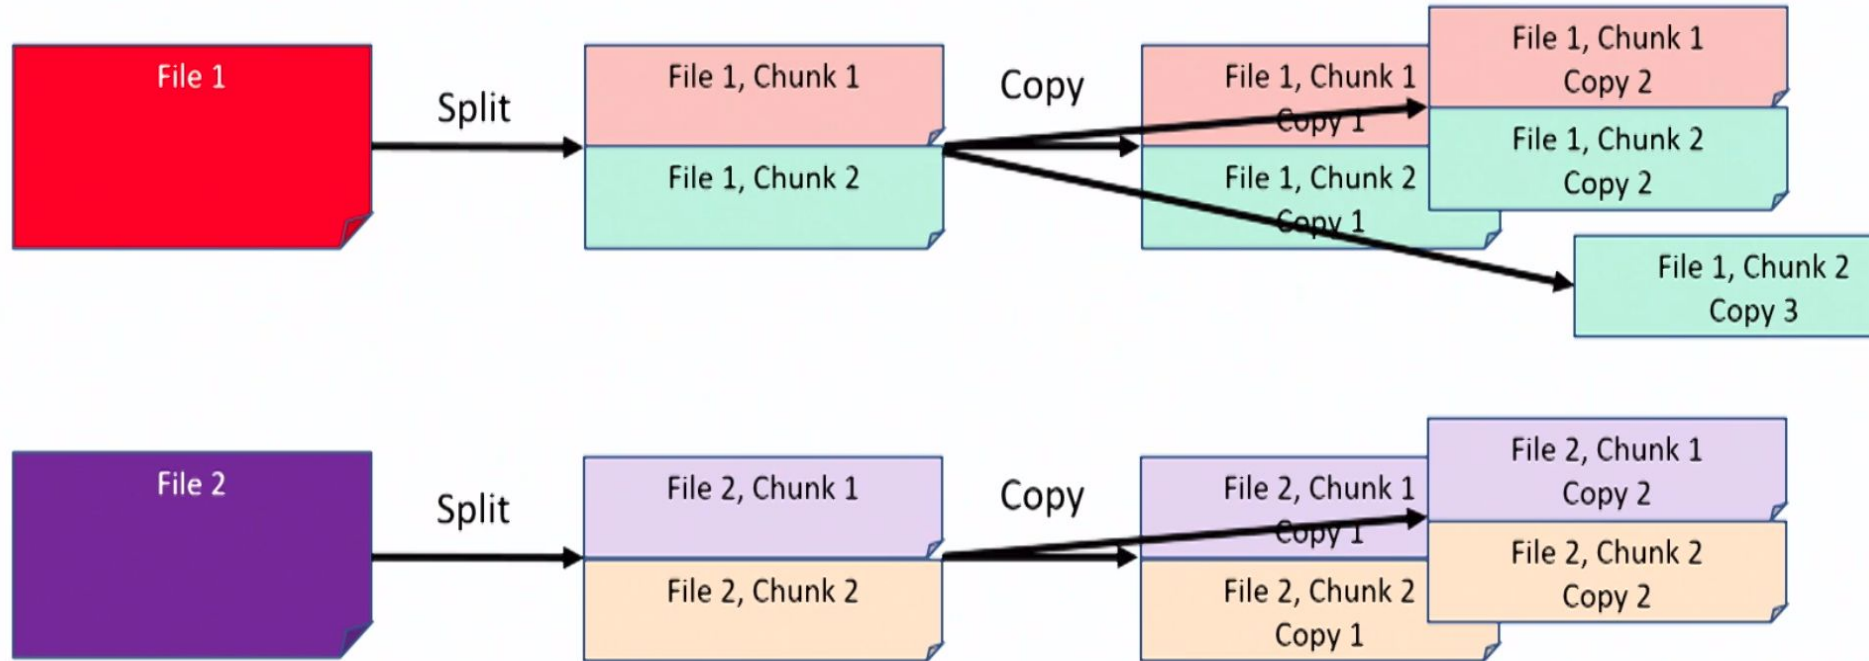
1. HDFS



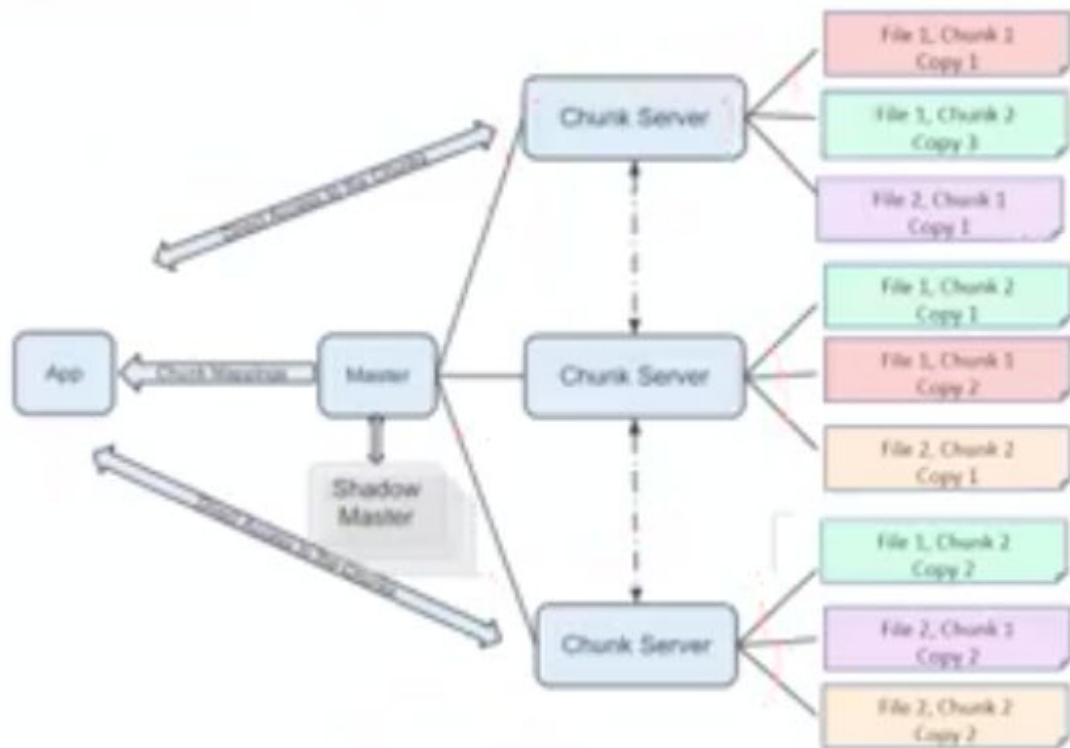


HDFS: Chunking files



HDFS: Chunking files



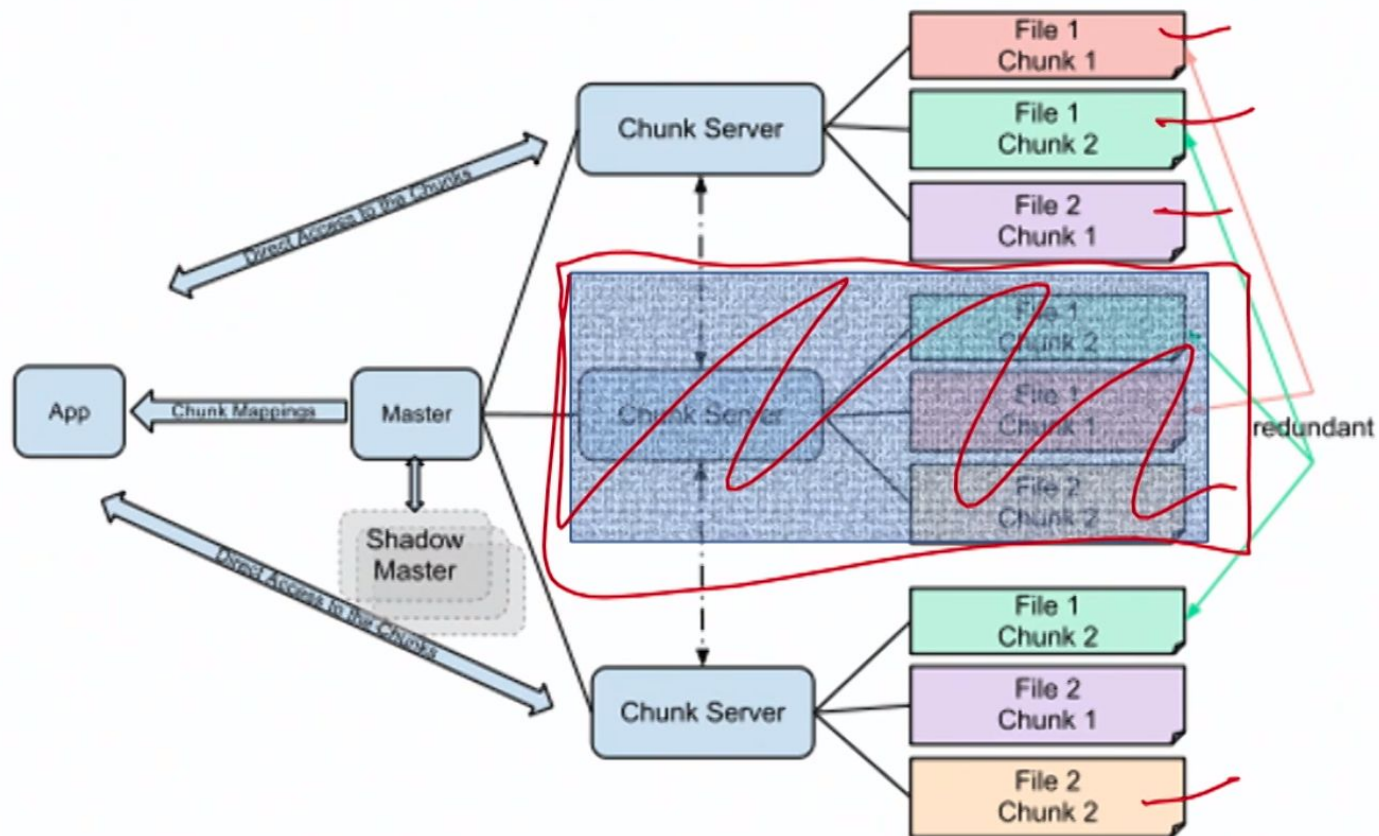
HDFS: Distributing Chunks



Properties of GFS/HDFS

- **Commodity Hardware** : Low cost per byte of storage
- **Locality** : data stored close to CPU
- **Redundancy** : can recover from server failures
- **Simple abstraction** : looks to user like standard file system (files,directories,etc) Chunk mechanism is hidden

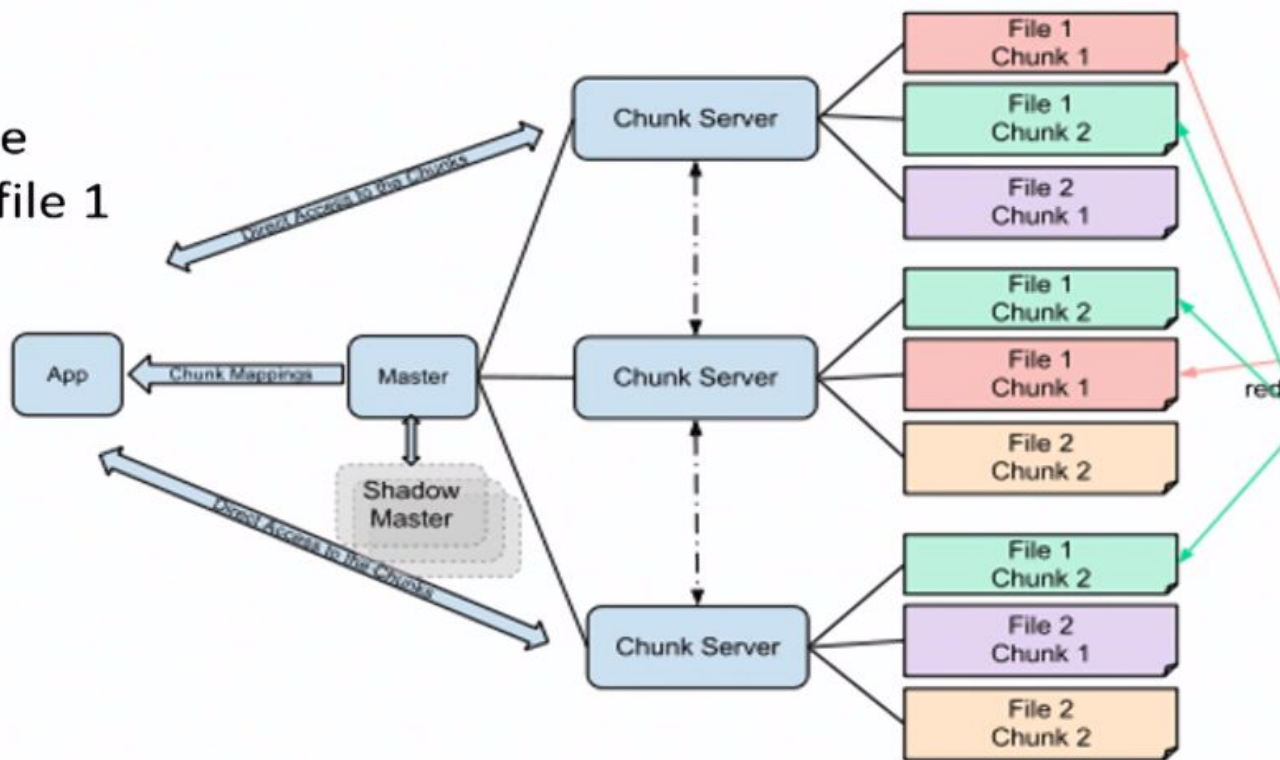
Redundancy



Locality

Task:

Sum all of the
elements in file 1




Map-Reduce

- HDFS is a **storage abstraction**
- Map-Reduce is a **computation abstraction** that works well with HDFS
- Allows programmer to specify parallel computation without knowing how hardware is organized


Summary

- Big data analysis is performed on large clusters of commodity computers
- HDFS (Hadoop file system) breakdown files to chunks, make copies, distribute randomly
- Hadoop Map-Reduce : a computation abstraction that works well with HDFS


2. Map-Reduce



Data Latency



Latencies



With big data, most of the latency is memory latency (1,2,4), not computation (3)

- Major source of latency in data is reading and writing into storage
- Different types of storage offer different latency, capacity and price
- Big Data analytics resolves around methods for organizing storage and computation in ways that maximize speed while minimizing cost.

Cache computations

How to use computers
clusters to efficiently
process large amounts
of data?

Reduce : compute the sum

- list L = [3,1,5,7]
- find the sum(16)

Traditional

```
## Use Builtin  
sum(L)  
  
## for loop  
s=0  
for i in L:  
    s+=i
```


Map-Reduce

```
reduce(lambda (x,y): x+y, L)
```


Order :

- Tradition : compute from first to last **in order**
- Map-Reduce : computation **order is not specified**

For loop order



Parallel order



Map + Reduce

- list L = [3,1,5,7]
- compute the sum of the squares

Traditional

```
## For Loop
s=0
for i in L:
    s+= i*i
## List comprehension
sum([i*i for i in L])
```

Map-Reduce

```
reduce(lambda x,y:x+y, \
        map(lambda i:i*i,L))
```

Order :

- Tradition : compute from first to last in **order**
- Map-Reduce : computation **order is not specified**

Execution :


- Tradition : immediate execution
 - Map-Reduce : execution plan
-

Map, Reduce operations should not depend on :

- Order of items in the list (**commutativity**)
- Order of operations (**associativity**)


4. RDD

Spark software architecture




- Driver - runs on master, and execute the main of the program
- Cluster Master - manage the computational resources
- Workers - manage each single core,
 - manage partitions and executors
- Each RDD is partitioned among the workers
- Executors - execute tasks on their partition

RDD Processing



- **Transformation** -> computation of RDD from the previous one
- **Action** -> trigger the computation
- **Lineage** -> one RDD is computed from another computed from another...
- RDDs, on default, are not materialized
- They do materialize if cached or otherwise persisted



- RDD graph -> execution plan defined by programmer, computation in head node
- Stage1 -> set of operations that can be done before materialization. You can compute quite a lot of RDD's because you don't need to actually store them
- Stage 2 -> data materialized

Transformations

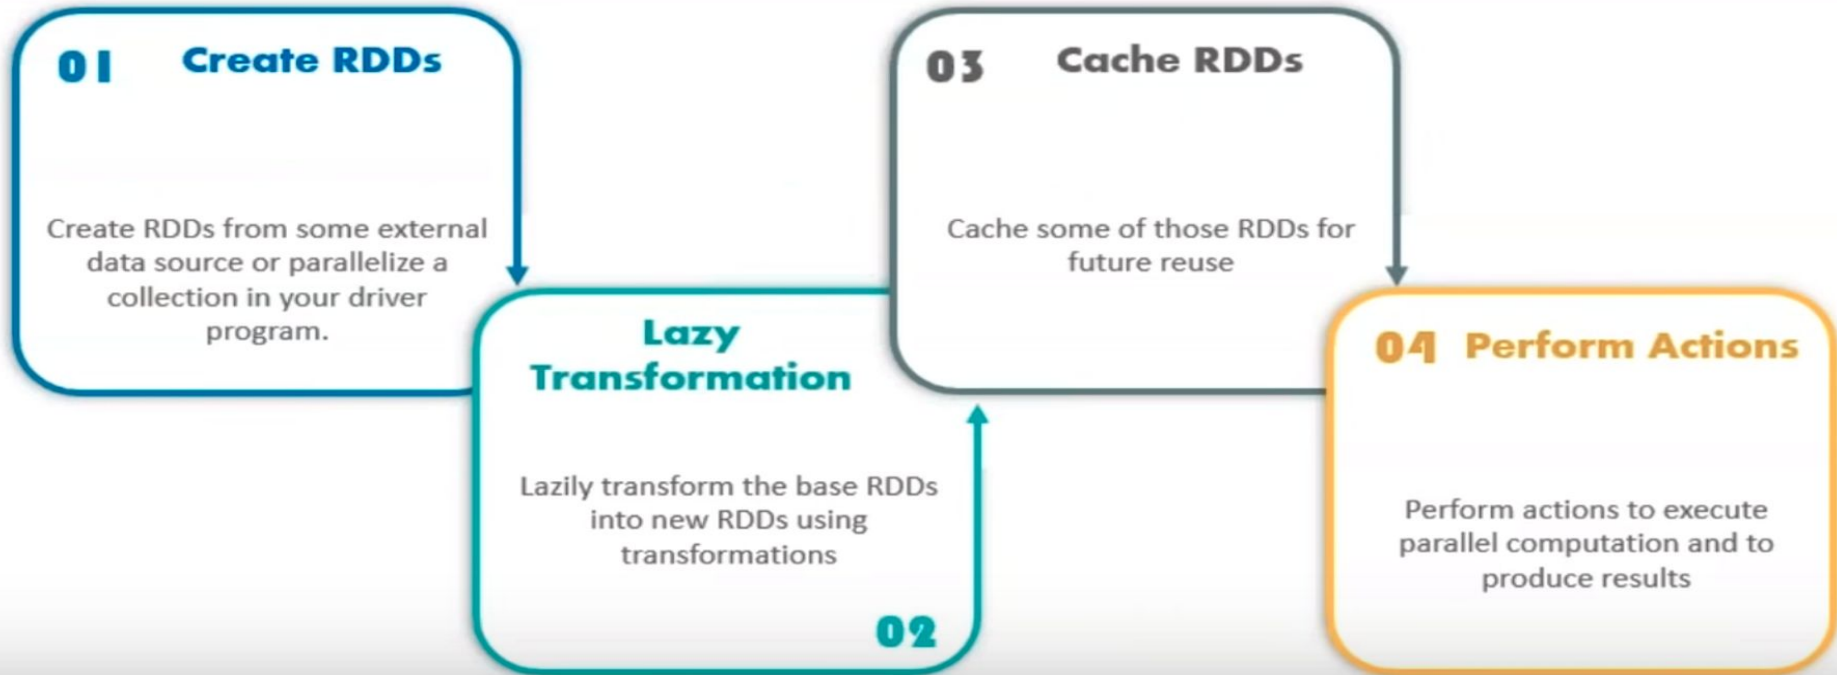
- map
- flatMap
- filter
- distinct
- reduceByKey
- mapPartitions
- sortBy

Actions

- collect
 - collectAsMap
 - reduce
 - countByKey/countByValue
 - take
 - first
-


Life Cycle of spark program

To work on this immutable data, you need to create a new one via Transformations and actions




Notebooki

Spark DataFrame




- **Distributed collection of data organized into *named columns***
- Conceptually equivalent to a table in **relational DB** or a **data frame in Pandas**
- API available in Scala, Java, Python, and R

UNSTRUCTURED



SEMI-STRUCTURED



STRUCTURED








More Flexible

///

Better Storage and Performance

DataFrame Transformations



-  Remove columns or rows
-  Transform a row into a column or a column into a row
-  Add rows or columns
-  Sort data by values in rows

Notebooki

Dziękuję za uwagę ;)

